

AMIR HAMZA KHAN

New Delhi, India | (+91) 7541896866 | amirhamzakhan2001@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

CAREER SUMMARY

AI/ML Engineer with hands-on experience building production-grade AI systems — including multi-tenant voice AI platforms (STT, LLM, TTS), Retrieval-Augmented Generation (RAG) pipelines, and no-code workflow automation tools and Agentic AI. Proficient in Python, FastAPI, LangChain, LangGraph, and Agentic AI/LLM application development, with a track record of shipping end-to-end products from architecture through deployment.

TECHNICAL SKILLS

Languages & Tools: Python, SQL, Git, Docker, MLflow, FastAPI, Jupyter Notebook, VS Code

ML & Deep Learning: PyTorch, Scikit-learn, Model Training, Transfer Learning, Feature Engineering, Hyperparameter Tuning, Model Evaluation

NLP & Generative AI: LangChain, LangGraph, Agentic AI, Multi-Agent Systems, RAG, Prompt Engineering, LLM Applications, Text Embeddings, Vector Search, Fine-Tuning, MLOps, LLMOps, STT, TTS, Claude

Data & Databases: Pandas, NumPy, Matplotlib, Seaborn, PostgreSQL, MongoDB, Vector Databases (FAISS, Qdrant), Power BI

PROFESSIONAL EXPERIENCE

AI Engineer Intern — Eyas Ventures, Noida

Feb 2026 – Jul 2026

Voice Agent Platform [\[link\]](#)

- Architected and built Voxa, a multi-tenant voice AI agent with configurable Speech-to-Text (STT), Large Language Model (LLM), and Text-to-Speech (TTS) pipelines per tenant, including voice cloning via ElevenLabs, Cartesia, and AssemblyAI, achieving ~1000ms end-to-end response latency.
- Developed the agent creation workflow — dynamic prompt generation, a Retrieval-Augmented Generation (RAG) knowledge base for company-specific responses, calendar integration, and a credit-based usage/billing system.
- Integrated 5+ external voice AI platforms (including Ringg, Retell, and Vapi) via REST APIs for cross-platform orchestration, and built 4-channel outbound campaign engines spanning call, SMS, email, and WhatsApp.
- Engineered end-to-end call data infrastructure — automated recordings, transcripts, and real-time analytics dashboards — giving businesses full visibility into agent and campaign performance.

Workflow Automation & CRM Integration

- Co-developed a no-code workflow automation builder (2-person team, later owned and extended independently) supporting conditional logic, scheduled and event-based triggers, webhooks, and multi-step automation chains.
- Integrated 10+ third-party applications and CRMs — including Gmail, Google Calendar, Google Sheets, WhatsApp, Telegram, OpenAI, and Claude — enabling businesses to automate cross-platform workflows without writing code.

Embeddable Website Widget

- Developed a lightweight, one-line script-tag embeddable widget enabling businesses to add AI-powered text and voice query support directly to their website, backed by the platform's RAG knowledge base.
- Built an admin-facing configuration layer allowing non-technical users to customize widget branding, appearance, and behavior without writing code.

AI Model Evaluation Researcher — Outlier AI, Remote

Oct 2025 – Feb 2026

- Evaluated Large Language Model (LLM) outputs against structured rubrics covering factual accuracy, hallucination detection, faithfulness, instruction-following, and safety compliance, delivering actionable feedback to improve model alignment.
- Identified recurring failure patterns and edge-case weaknesses across large-scale assessment tasks, contributing to measurable improvements in model response reliability.

PROJECTS

JMI University AI Chatbot & Query Management System

- Developing the AI chatbot layer for a university-wide query management platform covering 50+ departments, serving students, faculty, and prospective students, using a Retrieval-Augmented Generation (RAG) architecture powered by Qwen LLMs with ~1500ms response latency.
- Building backend services and chatbot orchestration in Python and FastAPI, and contributing to a multi-level, role-based admin portal for managing notices, circulars, documents, and chatbot knowledge sources.

NLP-Based Email Categorization System [\[Link\]](#)

- Built an end-to-end NLP pipeline processing 60,000+ Gmail emails using transformer-based embeddings and PCA-based dimensionality reduction for feature compression.
- Applied 4-level hierarchical clustering (Bisecting K-Means) to organize emails into 14 LLM-labeled clusters, then trained MLP classifiers for real-time email categorization — cutting manual sorting effort by an estimated 80–90%.

Cognivo — App & Web Platform for Dementia Patients [\[Link\]](#)

- Independently designed and developed Cognivo, a cross-platform care application (React Native, Android) and companion caregiver website with a FastAPI backend, to support dementia patients' daily independence.
- Implemented real-time face recognition (identifying familiar individuals in ~3–4 seconds) and a tap-to-ask computer vision feature for recognizing objects, medicines, and pets, plus scheduled reminders for patients.
- Built live location sharing, caregiver alerts for unrecognized visitors, two-way voice messaging, and custom voice cloning so the app delivers reminders in a familiar family member's voice.
- Self-taught mobile app development using AI-assisted (Claude) development workflows to independently ship a full-stack, multi-platform working prototype.

EDUCATION

M.Sc. Artificial Intelligence and Machine Learning *2024 – 2026*

Jamia Millia Islamia, New Delhi | CGPA: 9.49 | Ranked 1st in Class

B.Sc. (Hons.) Applied Mathematics *2019 – 2022*

Jamia Millia Islamia, New Delhi | CGPA: 9.53 | Ranked 2nd in Class

CERTIFICATIONS & ACHIEVEMENTS

- IBM AI Agent Architect Certificate
- Amazon ML Challenge 2025 — Built a multi-modal price prediction pipeline (PyTorch, Open CLIP, Sentence Transformers); SMAPE 54.6, ranked in top 25%.